

Sémantique des Langages de Programmation (SemLP)

DM : Cooperative Concurrency

The language

We consider a simple extension to the monadic *call-by-value* λ -calculus of the lecture notes (Chapter 14) with :

1. Natural values,
2. Boolean values,
3. String values,
4. an *if-then-else* conditional expression,
5. references Chapter 17, and
6. two concurrency constructs : **fork** and **yield**.

The syntax of this extended language, which we shall call concurrent- λ , is given in Figure 1.¹ The symbol \oplus represents a binary natural operator, \otimes represents a natural binary comparator, and \odot represents a boolean binary operator. The special value $()$ is named *unit*, and it is the only value of type **Unit**.

x	\in	$\mathcal{V}ar$	(var. names)
n	\in	\mathbb{N}	(naturals)
b	\in	$\{\text{true}, \text{false}\}$	(booleans)
s	\in	$[\text{a} - \text{zA} - \text{Z0} - \text{9}]^*$	(strings)
p	\in	$\mathcal{R}ef$	(references)
V	$::=$	$\lambda x. M \mid n \mid b \mid s \mid p \mid ()$	(values)
M	$::=$	$x \mid V \mid @(M, M)$	(terms)
		$\mid M \oplus M \mid -M \mid M \otimes M \mid M \odot M \mid \neg M$	
		$\mid \text{ref } M \mid !M \mid M := M$	
		$\mid \text{let } x = M \text{ in } M$	
		$\mid \text{if } M \text{ then } M \text{ else } M$	
		$\mid \text{fork } M \mid \text{yield}$	(concurrency)
		$\mid \text{print } M$	(print)
E	$::=$	$@(E, M) \mid @(V, E)$	(ev. contexts)
		$\mid E \oplus M \mid V \oplus E \mid -E \mid \dots$	
		$\mid \text{ref } E \mid !E \mid E := M \mid V := E$	
		$\mid \text{let } x = E \text{ in } M \mid \text{if } E \text{ then } M \text{ else } M$	
		$\mid \text{print } E$	

FIGURE 1 – Syntax for concurrent- λ .

The semantics of most of the construct of this language are as given in the lecture notes. The **print** construct does is a side-effecting operation which shows a value on the terminal. The evaluation of a **print** term results in a unit “()”.

1. You are free to **add** types and operators as you see fit.

Concurrency

The semantics of the `fork` M construct is to create a new thread of execution, returning a unit value “()”, and continue executing until termination (or until the following `yield`). This means that the expression M can now be evaluated concurrently with the expression that created the new thread. However, unlike concurrent programming languages such as Java, the execution of a thread can only be voluntarily interrupted to give way for the execution of another thread. This is achieved by the execution of the `yield` command.

As an example consider the following program :

```
print 0; fork (print 1); fork (print 2); yield; print 3
```

can produce any of the results : 0123, 0321, 0231, etc. Where interpret sequence “;” to be implemented as an application : $M;N = (\lambda x.N)M$ where the variable x does not appear in N . We also assume that program termination has the effect of a `yield`. Otherwise we can assume that each `fork` is called with a command terminated by `yield`.

As can be seen from the example above, we assume that the choice of the thread to resume after a `yield` command is non-deterministic (including the yielding thread). However, a realistic implementation should strive to provide a fair scheduler that gives equal opportunity to all threads. Then for instance, under a Round Robin scheduler, we would expect the result 0123 for the program above.

Exercises

Exercise 1 : Implementing basic concurrent- λ

1. Give an operational semantics similar to that of Chapter 12 of the lecture notes to concurrent- λ . Notice that you will need the *heap* to give semantics to references², and you will need to record the set of threads currently executing.
2. Implement in your favorite programming language³ an interpreter for concurrent- λ except for the concurrency constructs `fork` and `yield`. The input to the interpreter is an expression in the *abstract syntax* of the language. You are not required to provide a parser.⁴
3. Implement the CPS transformation of Chapter 10 of the lecture notes for concurrent- λ .

Exercise 2 : Implementing concurrency

1. Implement the `fork` and `yield` constructs of concurrent- λ on top of the transformations given above. Provide a Round Robin scheduler for this implementation.
2. Provide a different scheduler of your choice and provide an example program where the Round Robin scheduler differs from the one provided.
3. Provide a test suite showing that your implementation of concurrent- λ is correct. In particular, consider different placements for `yield` commands to show changes in the behavior of programs.

2. The evaluation contexts are given. Identifying the redexes is left as part of the exercise.

3. OCaml?

4. Providing a parser can be considered for extra credit.

Exercice 3 : Synchronization : Bonus

Consider adding an additional command `wait M` which blocks the execution of the current thread until the expression M evaluates to true. We assume that the expression M is effect-free. Whenever the expression does not evaluate to true, control is released as in the case of a `yield` command.

Submission

Submission Date : TBD

Submission Format : You have to submit the source code with appropriate :

1. README.md file explaining :
 - how to compile : should be completely automated,
 - how to run a set of predefined (by you) tests,
 - all the command line options, and if needed, the source syntax and format for your interpreter.
2. a builder file (Makefile, sbt, maven, ...),
3. the source code with reasonable comments,
4. a document providing the solutions to exercise 1.1, and a reasonable explanation of the tests you conducted with the expected results,
5. if your implementation does not immediately compile you might be requested to submit a virtual machine (or docker container) with a running implementation of your project.